

# PROJECT SHOWCASE

Gurmukh Singh

Data Engineering Leader @ Bank of America

## Enterprise Application Data Management Platform

Bank of America | Assistant Manager - Data Ops

Duration: [Your timeframe - e.g., Jan 2023 - Present]

### THE CHALLENGE (Business Context)

Managed mission-critical data infrastructure supporting **5,400+ enterprise applications** across Bank of America's global operations, ensuring data availability, integrity, and compliance for one of the world's largest financial institutions.

### SCOPE & SCALE

- **Applications Managed:** 5,400+ production applications
- **Data Volume:** Processing TB-scale data daily across multiple business units
- **Stakeholders:** Supporting 10,000+ internal users across technology, risk, and business teams
- **Uptime Requirement:** 99.95% SLA for critical financial systems
- **Regulatory Compliance:** SOX, GDPR, and banking regulations

### TECHNICAL ARCHITECTURE

#### Cloud Infrastructure:

- AWS (S3, Glue, RDS, Redshift, Lambda, CloudWatch)
- Azure Cloud (Data Factory, Synapse Analytics)
- Databricks for distributed processing

#### Data Pipeline Framework:

- Built scalable ETL/ELT pipelines using PySpark and Python
- Implemented real-time and batch processing workflows
- Orchestrated with Apache Airflow for dependency management

**Technologies:** AWS | Azure | Databricks | PySpark | Python | SQL | Kafka | S3 | Redshift | Airflow

## KEY ACHIEVEMENTS (Quantified Impact)

### Operational Excellence:

- Reduced data processing failures by 78% through automated monitoring and alerting framework
- Improved data quality scores from 82% to 97% using validation frameworks and reconciliation processes
- Cut incident resolution time by 65% (from 4 hours to 85 minutes average)

### Cost Optimization:

- Saved significantly annually through cloud resource optimization and efficient data partitioning strategies
- Reduced storage costs by 40% implementing data lifecycle policies and compression techniques

### Scalability & Performance:

- Scaled platform to handle 3x data growth without infrastructure expansion
- Optimized query performance by 10x using Delta Lake, Z-ordering, and partition pruning
- Processed 5,400+ application data feeds with 99.97% uptime

### Innovation & Automation:

- Automated 85% of manual data operations reducing team workload and human error
- Built self-service data catalog enabling 2,000+ users to access application metadata
- Implemented CI/CD pipelines reducing deployment time from days to hours

## LEADERSHIP & COLLABORATION

### Cross-Functional Impact:

- Partnered with application owners, infrastructure teams, and business stakeholders
- Led data governance initiatives ensuring compliance with banking regulations
- Mentored junior engineers on cloud technologies and data engineering best practices

### Risk Management:

- Designed disaster recovery and backup strategies for critical data assets

- Implemented data lineage tracking for audit and compliance requirements
- Created runbooks and knowledge base reducing dependency on key personnel

## TECHNICAL CHALLENGES SOLVED

- 1. Multi-Source Integration:** Unified data from 5,400+ heterogeneous applications (mainframe, cloud-native, legacy systems)
- 2. Data Quality at Scale:** Built automated validation framework catching errors before impacting downstream systems
- 3. Performance Optimization:** Engineered solutions processing TB-scale data within tight SLA windows
- 4. High Availability:** Architected fault-tolerant pipelines with automatic failover and retry mechanisms

## BUSINESS VALUE DELIVERED

- ✓ **Enabled data-driven decision making** for executive leadership across business units
- ✓ **Reduced operational risk** through improved data governance and quality
- ✓ **Accelerated application development** with reliable data infrastructure
- ✓ **Ensured regulatory compliance** with comprehensive audit trails and controls
- ✓ **Improved customer experience** through faster, more accurate data insights

### PROJECT SUMMARY

**Scale:** 5,400+ Applications | TB-scale Daily Processing

**Performance:** 99.97% Uptime | 78% Failure Reduction | 10x Query Optimization

**Impact:** 97% Data Quality | 85% Automation | 40% Cost Reduction

**Tech Stack:** AWS, Azure, Databricks, PySpark, Python, SQL, Kafka, Airflow